

210801159

April 23, 2023

1 Glass Classification

Assignment submitted by : *Ritvik Prasad M*

Registration Number : **210801159**

1.0.1 Context and Content

Context

This is a Glass Identification Data Set from UCI. It contains 9 attributes.

The response is glass type(discrete 7 values)

[Kaggle Link](#)

[UCI dataset Link](#)

Content

Attribute Information

- RI: refractive index
- Na: Sodium (unit measurement: weight percent in corresponding oxide, as are attributes 4-10)
- Mg: Magnesium
- Al: Aluminum
- Si: Silicon
- K: Potassium
- Ca: Calcium
- Ba: Barium
- Fe: Iron
- Type of glass: (class attribute)
 - 1 building_windows_float_processed

- 2 building_windows_non_float_processed
- 3 vehicle_windows_float_processed
- 4 vehicle_windows_non_float_processed (none in this database)
- 5 containers
- 6 tableware
- 7 headlamps

1.0.2 Data Loading and Analysis

```
[1]: # importing essential libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[74]: # loading the dataset
df = pd.read_csv("./glass.csv")
df.head()
```

```
[74]:
```

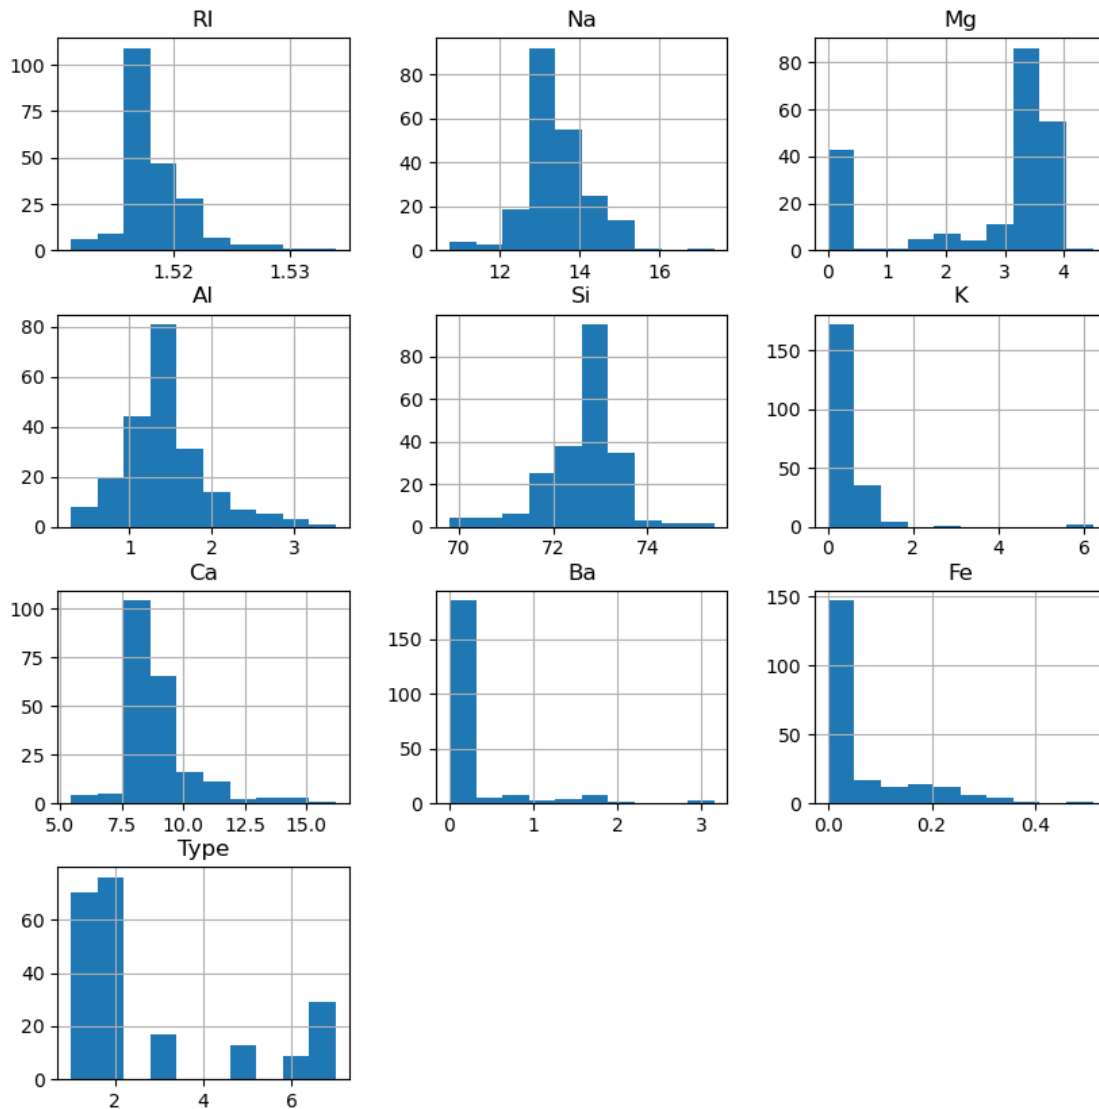
	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
0	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.0	0.0	1
1	1.51761	13.89	3.60	1.36	72.73	0.48	7.83	0.0	0.0	1
2	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0.0	0.0	1
3	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0.0	0.0	1
4	1.51742	13.27	3.62	1.24	73.08	0.55	8.07	0.0	0.0	1

```
[3]: #getting information about the data
df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 214 entries, 0 to 213
Data columns (total 10 columns):
#   Column  Non-Null Count  Dtype
---  -
0    RI      214 non-null      float64
1    Na      214 non-null      float64
2    Mg      214 non-null      float64
3    Al      214 non-null      float64
4    Si      214 non-null      float64
5    K       214 non-null      float64
6    Ca      214 non-null      float64
7    Ba      214 non-null      float64
8    Fe      214 non-null      float64
9    Type    214 non-null      int64
```

```
dtypes: float64(9), int64(1)
memory usage: 16.8 KB
```

```
[4]: # understanding the correlation between the data
df.hist(figsize=(10,10));
```



```
[5]: # analyzing the number and types of outputs
sns.countplot(x=df["Type"], width=0.7);
```



```
[6]: # finding the total number of possible outcomes
df.Type.unique(), len(df.Type.unique())
```

```
[6]: (array([1, 2, 3, 5, 6, 7]), 6)
```

Inference

As the data outcomes are discrete, we would use a multiclass classification algorithm such as non-binary multiclass classification.

For this data, I will use **Random Forest Classifier** Algorithm.

1.0.3 Model Creation and Training

```
[7]: # splitting the data into labels and outcomes
X = df.iloc[:, :-1]
y = df.iloc[:, -1]
X.shape , y.shape
```

```
[7]: ((214, 9), (214,))
```

```
[8]: # splitting the data into train and test set with test size of 0.2 as the
      ↪availability of data is less
      from sklearn.model_selection import train_test_split
      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)
```

```
[9]: # importing necessary libraries for model creation
      from sklearn.pipeline import Pipeline
      from sklearn.model_selection import GridSearchCV
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.preprocessing import StandardScaler
```

```
[10]: # creating the pipeline for the model and getting the pipeline parameters
      pipe = Pipeline([
          ("scaler", StandardScaler()),
          ("randomForest", RandomForestClassifier())
      ])
      pipe.get_params()
```

```
[10]: {'memory': None,
      'steps': [('scaler', StandardScaler()),
                ('randomForest', RandomForestClassifier())],
      'verbose': False,
      'scaler': StandardScaler(),
      'randomForest': RandomForestClassifier(),
      'scaler__copy': True,
      'scaler__with_mean': True,
      'scaler__with_std': True,
      'randomForest__bootstrap': True,
      'randomForest__ccp_alpha': 0.0,
      'randomForest__class_weight': None,
      'randomForest__criterion': 'gini',
      'randomForest__max_depth': None,
      'randomForest__max_features': 'sqrt',
      'randomForest__max_leaf_nodes': None,
      'randomForest__max_samples': None,
      'randomForest__min_impurity_decrease': 0.0,
      'randomForest__min_samples_leaf': 1,
      'randomForest__min_samples_split': 2,
      'randomForest__min_weight_fraction_leaf': 0.0,
      'randomForest__n_estimators': 100,
      'randomForest__n_jobs': None,
      'randomForest__oob_score': False,
      'randomForest__random_state': None,
      'randomForest__verbose': 0,
      'randomForest__warm_start': False}
```

```
[11]: # creating a cross validation estimator grid with estimators and maximum depth,
      ↪ as grid parameters
```

```
estimator = GridSearchCV(estimator=pipe, param_grid={
    'randomForest__n_estimators': [i for i in range(50,150,5)],
    'randomForest__max_depth': [i for i in range(1,11)]
}, scoring='accuracy', cv=3)
```

```
[52]: # fitting the estimator
estimator.fit(X_train, y_train)
```

```
[52]: GridSearchCV(cv=3,
                  estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                             ('randomForest',
                                              RandomForestClassifier())]),
                  param_grid={'randomForest__max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9,
                                                           10],
                              'randomForest__n_estimators': [50, 55, 60, 65, 70, 75,
                                                             80, 85, 90, 95, 100,
                                                             105, 110, 115, 120, 125,
                                                             130, 135, 140, 145]},
                  scoring='accuracy')
```

```
[53]: # storing the results of the estimator in a dataframe
result_df = pd.DataFrame(estimator.cv_results_)
```

```
[54]: # getting the parameters for the best optimum model
print(estimator.best_index_)
print(estimator.best_params_)
print(estimator.best_score_)
```

```
125
{'randomForest__max_depth': 7, 'randomForest__n_estimators': 75}
0.7543859649122807
```

```
[55]: # getting the details about best optimum model from result database
result_df.iloc[150]
```

```
[55]: mean_fit_time
0.122424
std_fit_time
0.002658
mean_score_time
0.010924
std_score_time
0.001117
param_randomForest__max_depth
```

```

8
param_randomForest__n_estimators
100
params                                {'randomForest__max_depth': 8,
'randomForest__...
split0_test_score
0.578947
split1_test_score
0.842105
split2_test_score
0.701754
mean_test_score
0.707602
std_test_score
0.107513
rank_test_score
93
Name: 150, dtype: object

```

```

[56]: # creating an improvised estimator with parameters close to best estimate of
      ↪previous estimator
est_depth = estimator.best_params_['randomForest__max_depth']
est_estimators = estimator.best_params_['randomForest__n_estimators']
delta1 = 5
delta2 = 2
estimator_improvised = GridSearchCV(estimator=pipe, param_grid={
    'randomForest__n_estimators': [i for i in
      ↪range(est_estimators-delta1,est_estimators+delta1+1)],
    'randomForest__max_depth': [i for i in
      ↪range(est_depth-delta2,est_depth+delta2+1)],
}, scoring='accuracy', cv=3)

```

```

[57]: # fitting the improvised estimator
estimator_improvised.fit(X_train, y_train)

```

```

[57]: GridSearchCV(cv=3,
                  estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                             ('randomForest',
                                              RandomForestClassifier())]),
                  param_grid={'randomForest__max_depth': [5, 6, 7, 8, 9],
                              'randomForest__n_estimators': [70, 71, 72, 73, 74, 75,
                                                              76, 77, 78, 79, 80]},
                  scoring='accuracy')

```

```

[58]: # storing the results of improvised estimator in a dataframe
result2_df = pd.DataFrame(estimator.cv_results_)

```

```
[59]: # getting the parameters for the best optimum model from the improvised_
      ↪ estimator
      print(estimator_improvised.best_index_)
      print(estimator_improvised.best_params_)
      print(estimator_improvised.best_score_)
```

```
27
{'randomForest__max_depth': 7, 'randomForest__n_estimators': 75}
0.7485380116959064
```

```
[60]: # getting the details about best optimum model from improvised result database
      result2_df.iloc[181]
```

```
[60]: mean_fit_time
      0.071356
      std_fit_time
      0.000997
      mean_score_time
      0.008694
      std_score_time
      0.000641
      param_randomForest__max_depth
      10
      param_randomForest__n_estimators
      55
      params                                {'randomForest__max_depth': 10,
      'randomForest_...
      split0_test_score
      0.596491
      split1_test_score
      0.894737
      split2_test_score
      0.684211
      mean_test_score
      0.725146
      std_test_score
      0.125152
      rank_test_score
      30
      Name: 181, dtype: object
```

```
[67]: # getting the top 2 estimators
      model1 = estimator.best_estimator_
      model2 = estimator_improvised.best_estimator_
```


1.0.4 Model Evaluation

```
[68]: # comparing the top two models

from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix

preds1 = model1.predict(X_test)
preds_train1 = model1.predict(X_train)
accuracy_test1 = accuracy_score(y_true=y_test, y_pred=preds1)
accuracy_train1 = accuracy_score(y_true=y_train, y_pred=preds_train1)
cmatrix1 = confusion_matrix(y_true=y_test, y_pred=preds1)
print(f'Model 1 Metrics')
print(f'accuracy (test)          : {accuracy_test1}\n\tin precentage :
      ↪\t{accuracy_test1*100} %\n')
print(f'accuracy (train)         : {accuracy_train1}\n\tin precentage :
      ↪\t{accuracy_train1*100} %\n')
print(f'confusion matrix (for 6 outcomes):\n{cmatrix1}')

preds2 = model2.predict(X_test)
preds_train2 = model2.predict(X_train)
accuracy_test2 = accuracy_score(y_true=y_test, y_pred=preds2)
accuracy_train2 = accuracy_score(y_true=y_train, y_pred=preds_train2)
cmatrix2 = confusion_matrix(y_true=y_test, y_pred=preds2)
print(f'\nModel 2 Metrics')
print(f'accuracy (test)          : {accuracy_test2}\n\tin precentage :
      ↪\t{accuracy_test2*100} %\n')
print(f'accuracy (train)         : {accuracy_train2}\n\tin precentage :
      ↪\t{accuracy_train2*100} %\n')
print(f'confusion matrix (for 6 outcomes):\n{cmatrix2}')
```

Model 1 Metrics

accuracy (test) : 0.9069767441860465
 in precentage : 90.69767441860465 %

accuracy (train) : 0.9766081871345029
 in precentage : 97.6608187134503 %

confusion matrix (for 6 outcomes):

```
[[11  0  0  0  0  0]
 [ 3 11  0  0  0  0]
 [ 1  0  2  0  0  0]
 [ 0  0  0  4  0  0]
 [ 0  0  0  0  3  0]
 [ 0  0  0  0  0  8]]
```

Model 2 Metrics

accuracy (test) : 0.813953488372093

```
in precentage : 81.3953488372093 %  
  
accuracy (train)      : 0.9766081871345029  
in precentage : 97.6608187134503 %
```

confusion matrix (for 6 outcomes):

```
[[11  0  0  0  0  0]  
 [ 3 10  0  0  0  1]  
 [ 1  1  1  0  0  0]  
 [ 0  2  0  2  0  0]  
 [ 0  0  0  0  3  0]  
 [ 0  0  0  0  0  8]]
```

```
[75]: print(f'Test delta  : {(accuracy_test1-accuracy_test2)*100}%')  
      print(f'Train delta : {(accuracy_train1-accuracy_train2)*100}%')
```

```
Test delta  : 9.302325581395344%  
Train delta : 0.0%
```

We can see the metrics as:

- Model1
 - Test Accuracy : 90.69%
 - Train Accuracy : 97.66%
- Model2
 - Test Accuracy : 81.39%
 - Train Accuracy : 97.66%

We can see that *model1* perfoms better on train set than test set by a margin of **9.30%** and *model1* perfoms better on test set than train set by a margin of **0.00%**.

Hence, from this observation, we can say that model1 outperforms model2.

Therefore model1 is the most optimum model.

1.0.5 Exporting the Model

```
[71]: # exporting the model in .joblib format  
      from joblib import dump  
      dump(model1, "glass_pred.joblib")
```

```
[71]: ['glass_pred.joblib']
```

```
[73]: # loading the model and finding its accuracy  
  
      from joblib import load  
      from sklearn.metrics import accuracy_score
```

```
load_model = load("./glass_pred.joblib")

pred_load = load_model.predict(X_test)
accuracy_load_model = accuracy_score(y_true=y_test,y_pred=pred_load)
print(f'Accuracy of the model loaded : {accuracy_load_model*100} %')
```

Accuracy of the model loaded : 90.69767441860465 %

1.0.6 Conclusion

Hence, a **Random Forest Classifier** algorithm is implemented on the given data and a Machine Learning model with **90.698% accuracy** is obtained.