

glass-pred

April 22, 2023

1 Glass Classification

Assignment submitted by : *Ritvik Prasad M*

Registration Number : **210801159**

Context

This is a Glass Identification Data Set from UCI. It contains 9 attributes.

The response is glass type(discrete 7 values)

[Kaggle Link](#)

[UCI dataset Link](#)

Content

Attribute Information - RI: refractive index - Na: Sodium (unit measurement: weight percent in corresponding oxide, as are attributes 4-10) - Mg: Magnesium - Al: Aluminum - Si: Silicon - K: Potassium - Ca: Calcium - Ba: Barium - Fe: Iron - Type of glass: (class attribute)
- 1 building_windows_float_processed - 2 building_windows_non_float_processed - 3 vehicle_windows_float_processed - 4 vehicle_windows_non_float_processed (none in this database)
- 5 containers - 6 tableware - 7 headlamps

1.0.1 Data Loading and Analysis

```
[107]: # importing essential libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[108]: # loading the model
df = pd.read_csv("./glass.csv")
df.head()
```

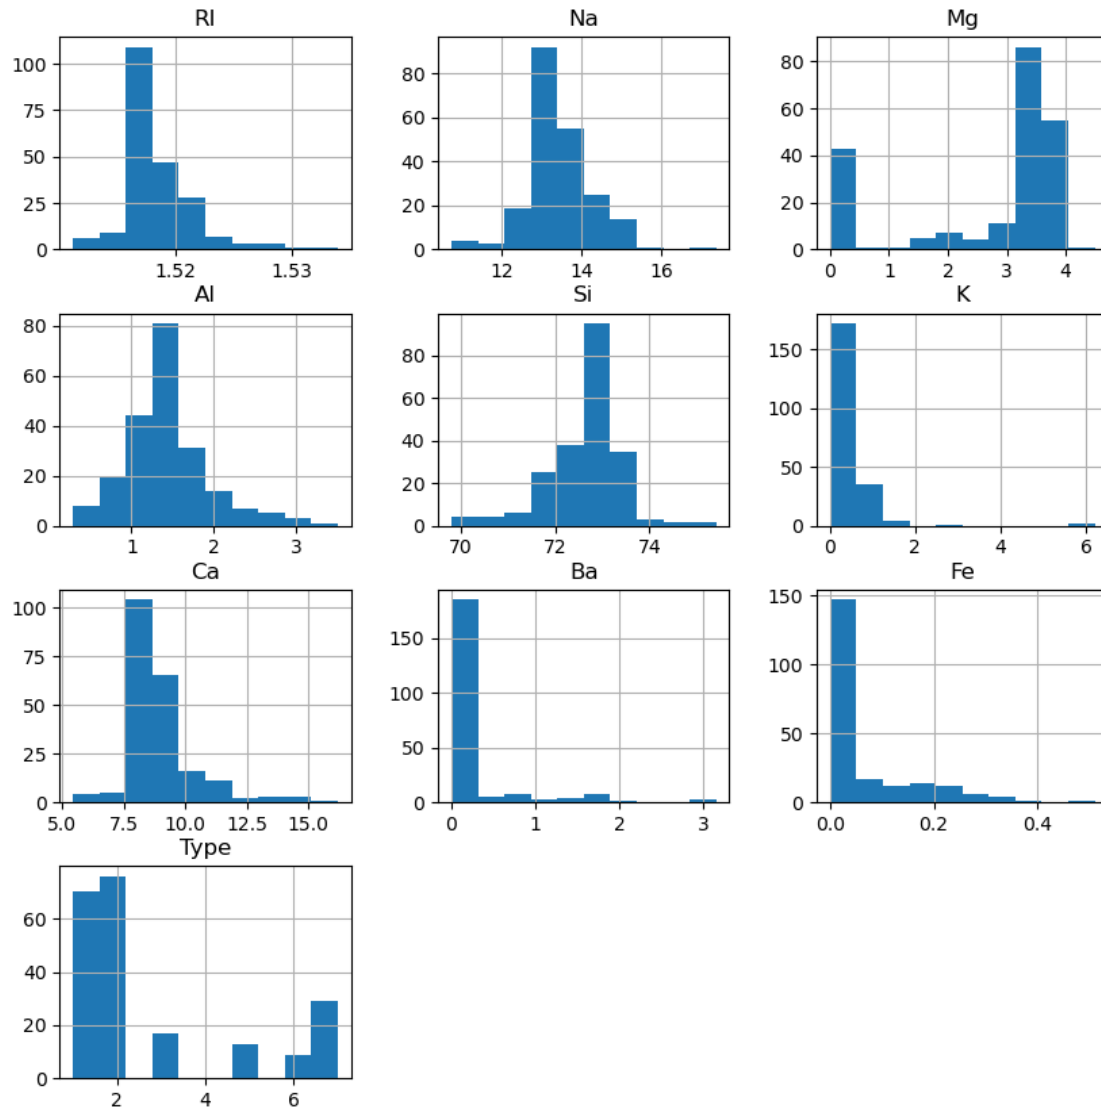
```
[108]:
```

	RI	Na	Mg	Al	Si	K	Ca	Ba	Fe	Type
0	1.52101	13.64	4.49	1.10	71.78	0.06	8.75	0.0	0.0	1
1	1.51761	13.89	3.60	1.36	72.73	0.48	7.83	0.0	0.0	1
2	1.51618	13.53	3.55	1.54	72.99	0.39	7.78	0.0	0.0	1
3	1.51766	13.21	3.69	1.29	72.61	0.57	8.22	0.0	0.0	1
4	1.51742	13.27	3.62	1.24	73.08	0.55	8.07	0.0	0.0	1


```
[109]: #getting information about the data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 214 entries, 0 to 213
Data columns (total 10 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    RI      214 non-null    float64
 1    Na      214 non-null    float64
 2    Mg      214 non-null    float64
 3    Al      214 non-null    float64
 4    Si      214 non-null    float64
 5    K       214 non-null    float64
 6    Ca      214 non-null    float64
 7    Ba      214 non-null    float64
 8    Fe      214 non-null    float64
 9    Type    214 non-null    int64
dtypes: float64(9), int64(1)
memory usage: 16.8 KB
```

```
[110]: # understanding the correlation between the data
df.hist(figsize=(10,10));
```



```
[111]: # analyzing the number and types of outputs
sns.countplot(x=df["Type"], width=0.7);
```



```
[112]: # finding the total number of possible outcomes
df.Type.unique(), len(df.Type.unique())
```

```
[112]: (array([1, 2, 3, 5, 6, 7]), 6)
```

Inference

As the data outcomes are discrete, we would use a multiclass classification algorithm such as non-binary multiclass classification.

For this data, I will use **Random Forest Classifier** Algorithm.

1.0.2 Model Creation and Training

```
[113]: # splitting the data into labels and outcomes
X = df.iloc[:, :-1]
y = df.iloc[:, -1]
X.shape , y.shape
```

```
[113]: ((214, 9), (214,))
```

```
[114]: # splitting the data into train and test set with test size of 0.2 as the
        ↪availability of data is less
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
        ↪random_state=42)
```

```
[115]: # importing necessary libraries for model creation
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
```

```
[116]: # creating the pipeline for the model and getting the pipeline parameters
pipe = Pipeline([
    ("scaler", StandardScaler()),
    ("randomForest", RandomForestClassifier())
])
pipe.get_params()
```

```
[116]: {'memory': None,
      'steps': [('scaler', StandardScaler()),
                ('randomForest', RandomForestClassifier())],
      'verbose': False,
      'scaler': StandardScaler(),
      'randomForest': RandomForestClassifier(),
      'scaler__copy': True,
      'scaler__with_mean': True,
      'scaler__with_std': True,
      'randomForest__bootstrap': True,
      'randomForest__ccp_alpha': 0.0,
      'randomForest__class_weight': None,
      'randomForest__criterion': 'gini',
      'randomForest__max_depth': None,
      'randomForest__max_features': 'sqrt',
      'randomForest__max_leaf_nodes': None,
      'randomForest__max_samples': None,
      'randomForest__min_impurity_decrease': 0.0,
      'randomForest__min_samples_leaf': 1,
      'randomForest__min_samples_split': 2,
      'randomForest__min_weight_fraction_leaf': 0.0,
      'randomForest__n_estimators': 100,
      'randomForest__n_jobs': None,
      'randomForest__oob_score': False,
      'randomForest__random_state': None,
      'randomForest__verbose': 0,
      'randomForest__warm_start': False}
```

```
[117]: # creating a cross validation estimator grid with estimators and maximum depth,
        ↪ as grid parameters
```

```
estimator = GridSearchCV(estimator=pipe, param_grid={
    'randomForest__n_estimators': [i for i in range(50,150,5)],
    'randomForest__max_depth': [i for i in range(1,11)]
}, scoring='accuracy', cv=3)
```

```
[118]: # fitting the estimator
estimator.fit(X_train, y_train)
```

```
[118]: GridSearchCV(cv=3,
                    estimator=Pipeline(steps=[('scaler', StandardScaler()),
                                                ('randomForest',
                                                 RandomForestClassifier())]),
                    param_grid={'randomForest__max_depth': [1, 2, 3, 4, 5, 6, 7, 8, 9,
                                                            10],
                                'randomForest__n_estimators': [50, 55, 60, 65, 70, 75,
                                                                80, 85, 90, 95, 100,
                                                                105, 110, 115, 120, 125,
                                                                130, 135, 140, 145]},
                    scoring='accuracy')
```

```
[119]: # storing the results of the estimator in a dataframe
result_df = pd.DataFrame(estimator.cv_results_)
```

```
[120]: # getting the parameters for the best optimum model
print(estimator.best_index_)
print(estimator.best_params_)
print(estimator.best_score_)
```

```
145
{'randomForest__max_depth': 8, 'randomForest__n_estimators': 75}
0.7543859649122807
```

```
[121]: # getting the details about best optimum model from result database
result_df.iloc[150]
```

```
[121]: mean_fit_time
0.135819
std_fit_time
0.00662
mean_score_time
0.009605
std_score_time
0.000944
param_randomForest__max_depth
```

```

8
param_randomForest__n_estimators
100
params                                {'randomForest__max_depth': 8,
'randomForest__...
split0_test_score
0.561404
split1_test_score
0.824561
split2_test_score
0.719298
mean_test_score
0.701754
std_test_score
0.108148
rank_test_score
111
Name: 150, dtype: object

```

```

[122]: # creating an improvised estimator with parameters close to best estimate of
↳previous estimator

```

```

estimator_improvised = GridSearchCV(estimator=pipe, param_grid={
    'randomForest__n_estimators':[i for i in range(50,61)],
    'randomForest__max_depth':[i for i in range(7,14)]
}, scoring='accuracy', cv=3)

```

```

[123]: # fitting the improvised estimator
estimator_improvised.fit(X_train, y_train)

```

```

[123]: GridSearchCV(cv=3,
    estimator=Pipeline(steps=[('scaler', StandardScaler()),
                              ('randomForest',
                               RandomForestClassifier())]),
    param_grid={'randomForest__max_depth': [7, 8, 9, 10, 11, 12, 13],
                'randomForest__n_estimators': [50, 51, 52, 53, 54, 55,
                                                56, 57, 58, 59, 60]},
    scoring='accuracy')

```

```

[124]: # storing the results of improvised estimator in a dataframe
result2_df = pd.DataFrame(estimator.cv_results_)

```

```

[125]: result2_df.to_csv("result2.csv")

```

```

[126]: # getting the parameters for the best optimum model from the improvised
↳estimator
print(estimator_improvised.best_index_)

```

```
print(estimator_improvised.best_params_)
print(estimator_improvised.best_score_)
```

```
5
{'randomForest__max_depth': 7, 'randomForest__n_estimators': 55}
0.7719298245614036
```

```
[127]: # getting the details about best optimum model from improvised result database
result2_df.iloc[181]
```

```
[127]: mean_fit_time
0.078363
std_fit_time
0.000689
mean_score_time
0.006414
std_score_time
0.001151
param_randomForest__max_depth
10
param_randomForest__n_estimators
55
params                                {'randomForest__max_depth': 10,
'randomForest__...
split0_test_score
0.596491
split1_test_score
0.842105
split2_test_score
0.77193
mean_test_score
0.736842
std_test_score
0.103295
rank_test_score
15
Name: 181, dtype: object
```

```
[128]: # getting the top 2 estimators
model1 = estimator.best_estimator_
model2 = estimator_improvised.best_estimator_
```


1.0.3 Model Evaluation

```
[143]: # comparing the top two models

from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
preds = model1.predict(X_test)
preds_train = model1.predict(X_train)
accuracy_test = accuracy_score(y_true=y_test, y_pred=preds)
accuracy_train = accuracy_score(y_true=y_train, y_pred=preds_train)
cmatrix = confusion_matrix(y_true=y_test, y_pred=preds)
print(f'Model 1 Metrics')
print(f'accuracy (test)          : {accuracy_test}\n\tin precentage :
      ↪\t{accuracy_test*100} %\n')
print(f'accuracy (train)         : {accuracy_train}\n\tin precentage :
      ↪\t{accuracy_train*100} %\n')
print(f'confusion matrix (for 6 outcomes):\n{cmatrix}')

preds = model2.predict(X_test)
preds_train = model2.predict(X_train)
accuracy_test = accuracy_score(y_true=y_test, y_pred=preds)
accuracy_train = accuracy_score(y_true=y_train, y_pred=preds_train)
cmatrix = confusion_matrix(y_true=y_test, y_pred=preds)
print(f'\nModel 2 Metrics')
print(f'accuracy (test)          : {accuracy_test}\n\tin precentage :
      ↪\t{accuracy_test*100} %\n')
print(f'accuracy (train)         : {accuracy_train}\n\tin precentage :
      ↪\t{accuracy_train*100} %\n')
print(f'confusion matrix (for 6 outcomes):\n{cmatrix}')
```

Model 1 Metrics

accuracy (test) : 0.8837209302325582
 in precentage : 88.37209302325581 %

accuracy (train) : 0.9883040935672515
 in precentage : 98.83040935672514 %

confusion matrix (for 6 outcomes):

```
[[11  0  0  0  0  0]
 [ 3 11  0  0  0  0]
 [ 1  0  2  0  0  0]
 [ 0  1  0  3  0  0]
 [ 0  0  0  0  3  0]
 [ 0  0  0  0  0  8]]
```

Model 2 Metrics

accuracy (test) : 0.8604651162790697
 in precentage : 86.04651162790698 %

```
accuracy (train)      : 1.0
    in precentage    : 100.0 %
```

confusion matrix (for 6 outcomes):

```
[[11  0  0  0  0  0]
 [ 1 11  1  0  0  1]
 [ 1  0  2  0  0  0]
 [ 0  2  0  2  0  0]
 [ 0  0  0  0  3  0]
 [ 0  0  0  0  0  8]]
```

We can see that *model2* perfoms better on train set than test set by a margin of **1.17%** and *model1* perfoms better on test set than train set by a margin of **2.325%**.

Hence, from this observation, we can say that *modell1* outperforms *model2*.

Therefore *modell1* is the most optimum model.

1.0.4 Exporting the Model

```
[145]: # exporting the model in .joblib format
from joblib import dump
dump(model1, "glass_pred.joblib")
```

```
[145]: ['glass_pred.joblib']
```

```
[146]: # loading the model and finding its accuracy

from joblib import load
from sklearn.metrics import accuracy_score

load_model = load("./glass_pred.joblib")

pred_load = load_model.predict(X_test)
accuracy_load_model = accuracy_score(y_true=y_test,y_pred=pred_load)
print(f'Accuracy of the model loed : {accuracy_load_model*100} %')
```

```
Accuracy of the model loed : 88.37209302325581 %
```

1.0.5 Conclusion

Hence, a **Random Forest Classifier** algorithm is implemented on the given data and a Machine Learning model with 88% accuracy is obtained.